

Automated Reasoning in Generating Exam Sheets for Automated Deduction

Petra Hozzová, Laura Kovács, and Jakob Rath

TU Wien, Austria

Abstract. Amid the COVID-19 pandemic, distance teaching became default in higher education, urging teachers and researchers to revise course materials into an accessible online content for a diverse audience. Probably one of the hardest challenges came with online assessments of course performance, for example by organizing online written exams. In this short paper we describe the online setting we organized for our master’s level course “Automated Deduction” in logic and computation at TU Wien. The algorithmic and rigorous reasoning developed within our course called for individual exam sheets focused on problem solving and deductive proofs; as such exam sheets using test grids were not a viable solution for written exams within our course. We believe the toolchain of automated reasoning tools we have developed for holding online written exams could be beneficial not only for other distance learning platforms, but also to researchers in automated reasoning, by providing our community with a large set of randomly generated benchmarks in SAT/SMT solving and first-order theorem proving.

1 Motivation

Amid the COVID-19 pandemic, higher education has moved to distance teaching. While online lecturing was relatively fast to implement via webinars, recordings, streaming and online communication channels, coming up with best practices to assess course performance was far from trivial. Even with very sophisticated technical infrastructure (use of which, on the other hand, would be unethical to require from course participants), avoiding collusion in a virtual environment is very hard to achieve, if possible at all. While work on online feedback generation has already been initiated, see e.g. [7,14], not much work on online examinations has so far emerged.

In this paper we present our approach to organizing online written exams, where the exam solutions require rigorous logical reasoning and proofs rather than using mechanized test grids. In particular, we were faced with the challenge of organizing online written exams for

our master’s level course “Automated Deduction” in logic and computation at TU Wien¹. This course introduces algorithmic techniques and fundamental results in automated reasoning, by focusing on specialised algorithms for reasoning in various fragments of first-order logics, such as propositional logic, combinations of ground theories, and full first-order logic with equality. As such, topics of the course cover theoretical and practical aspects of SAT/SMT solving [4,12,5] and first-order theorem proving using superposition reasoning [11,8].

We claim by no means that the framework we developed for online examination is optimal. Given the time constraints of examination periods, we aimed for an online exam setting that (i) reduces collusion among students and (ii) requires the same workload on each participant. The algorithmic reasoning developed within our course called for exam sheets focused on problem solving and deductive proofs; exam sheets using test grids were therefore not a viable solution for written exams within our course. We have therefore used and adapted the automated reasoning approaches introduced in our course to automate the generation of individual exam sheets for students enrolled in our course, by making sure that the exam tasks remain essentially the same in each generated exam sheet. As such, we have randomly generated individual exam problems on

- SAT solving, by imposing (mostly) syntactical constraints on randomly generated SAT formulas (Section 2.1);
- Satisfiability modulo theory (SMT) reasoning, by exploiting reasoning in a combination of theories and vary patterns of SMT problem templates (Section 3.1);
- First-order theorem proving, by adjusting simplification orderings in superposition reasoning and using redundancy elimination in first-order proving, both in the ground/quantifier-free and non-ground/quantified setting (Section 2.2 and Section 3.2).

For each of the SMT and first-order problems we generated, we used respective SMT and first-order solvers to perform an additional sanity check (Section 4). Our toolchain is available at <https://github.com/JakobR/exagen>.

¹ <https://tiss.tuwien.ac.at/course/courseDetails.xhtml?dsrid=2002&dsrid=601&courseNr=184774&semester=2020S>

We believe our framework could be beneficial not only for other distance learning platforms, but also to researchers in automated reasoning, by providing our community with a large set of randomly generated benchmarks in SAT/SMT solving and first-order theorem proving. While our proposal is specific to the formal aspects of automated reasoning, we note that our framework can be extended with further constraints to scale it to other courses in formal methods.

2 Random Problem Generation

We first describe our solution for generating automated reasoning benchmarks in a fully automated and random manner. We used this setting for generating exam problems on SAT solving and first-order theorem proving, by filtering out problem instances that are either too hard or too easy. Throughout this paper, we assume basic familiarity with standard first-order logic and refer to the literature [2,8] for further details.

2.1 Boolean Satisfiability (SAT)

In our exam problem on SAT solving, students were asked to (a) determine which atoms are of pure polarity in the formula, (b) compute a polarity-optimized clausal normal form [13], and (c) decide satisfiability of the computed CNF formula by applying the DPLL algorithm.

Randomly generating propositional formulas in a naive setting would lead to a huge variety of formulas, spanning both formulas for which the above questions are trivial to answer (e.g. clauses as propositional tautologies) and others requiring much more effort (e.g. arbitrary formulas using only \leftrightarrow). More work was thus needed to ensure comparable workload for solving exam sheets.

To this end, we identified several syntactical characteristics that the exam problems on SAT solving should exhibit, and filtered the generated formulas by these, as summarized partially below.

- (1) The SAT formula contains exactly seven logical connectives and exactly three different propositional variables.
- (2) There is at least one atom that appears with a pure polarity.

- (3) The connectives “ \leftrightarrow ”, “ \rightarrow ”, and “ \neg ” appear at least once, with “ \leftrightarrow ” appearing at most twice. At least one of “ \wedge ” and “ \vee ” appears.
- (4) The polarity-optimized clausal normal form results in a set of definitions, each of which is of the form $n \circ \varphi$ with $\circ \in \{\rightarrow, \leftarrow, \leftrightarrow\}$, a fresh propositional variable n , and a formula φ . We restrict the SAT formula such that at least two of the choices for \circ appear in the clausal normal form (CNF).
- (5) The SAT formula has at most 6 models.

While the combination of the above conditions might seem very restrictive, we note that there are 20 390 076 different SAT formulas satisfying the above criteria. Further, if we do not want to distinguish formulas that differ only by a permutation of atoms, 3 398 346 formulas remain.

2.2 Non-Ground Superposition with Redundancy

Moving beyond boolean satisfiability, we developed a random problem generator for first-order formulas with equality, in the setting of superposition-based first-order theorem proving with redundancy elimination [11, 8]. In this problem, a concrete inference² was given to the students, and their task was to (a) prove that the inference is sound and (b) that the inference is a simplification inference.

We recall that a simplification inference is an inference that removes clauses from the proof search space, whereas a generating inferences add new clauses to the search space [8]. In our work, we considered the simplification inference of *subsumption resolution*

$$\frac{A \vee C \quad \neg B \vee D}{D} \quad \text{or} \quad \frac{\neg A \vee C \quad B \vee D}{D} \quad (1)$$

where A, B are atoms and C, D are clauses such that for some substitution (or mgu) θ we have $A\theta \vee C\theta \subseteq B \vee D$, and hence the second premise $\neg B \vee D$ (or $B \vee D$) of (1) is redundant and can be deleted from the search space after applying (1) within proof search. We randomly generated first-order instances of the inference rule (1), as discussed next. Our setting could however be easily extended to other

² i.e., an instance of an inference rule as opposed to the rule itself

simplification inferences such as subsumption demodulation [6], and even generating inferences.

- (1) To randomly generate first-order terms and literals, we fixed a first-order signature consisting of predicate and function symbols and specified a set of logical variables. We controlled the shape of the generated terms by giving bounds on the *depth* of the term, that is the maximal nesting level of function calls (e.g., a constant symbol b has depth 0, while the term $g(f(x), d)$ has depth 2).
- (2) To obtain random instances of (1), we first generated non-ground clauses $C_2 := L_1 \vee L_2$ corresponding to an instance of the first premise of (1). To this end, we generated a random uninterpreted literal L_1 containing exactly one variable occurrence, and a random equality literal L_2 containing at least two occurrences of a different variable.
- (3) We next generated the clause $C_1 := \overline{L_1}\theta \vee L_2\theta \vee L_3$ as an instance of the second premise of (1) where θ is a randomly generated grounding substitution, L_3 is a randomly generated ground literal, and \overline{L} is the complementary³ literal to L .
- (4) We set $C_3 := L_3 \vee L_2\theta$ as an instance of the conclusion of (1), yielding thus the inference $\frac{C_1 \quad C_2}{C_3}$ as an instance of (1).

We found that with the concrete signature used for our exam, the discussed method can generate more than 10^{11} different instances of this inference.

3 Random Variation of Problem Templates

We now describe our framework for generating random quantifier-free first-order formulas with and without theories, that were used in the SMT reasoning and ground superposition proving tasks of our exam. For both of these tasks, we used quantifier-free first-order formula templates and implemented randomization over these templates by considering theory reasoning and simplification orderings.

3.1 Satisfiability Modulo Theories (SMT)

We considered first-order formula templates in the combined, quantifier-free theories of equality, arrays and linear integer arithmetic, corre-

³ i.e., $\overline{L} = \neg L$ and $\neg \overline{L} = L$.

sponding to the logic AUFLIA of SMT-LIB [1]. We aimed at generating SMT formulas over which reasoning in all three theories was needed, by exploiting the DPLL(T) framework [12] in combination with the Nelson-Oppen decision procedure [10].

With a naive random generation, it might however happen that, for example, array reasoning is actually not needed to derive (un)satisfiability of the generated SMT formula. We therefore constructed an SMT formula template and randomly introduced small perturbations in this template, so that the theory-specific reasoning in all generated SMT instances is different while reasoning in all theories is necessary. For doing so, we considered an SMT template with two constants of integer sorts and replaced an integer-sorted constant symbol c by integer-sorted terms $c + i$, where $i \in \{-3, -2, \dots, 3\}$ is chosen randomly. We flattened nested arithmetic terms such as $(c + i) + j$ to $c + k$, where i, j, k are integers and $k = i + j$. As a result, we generated 49 different SMT problems; we illustrate one such formula, together with its reasoning tasks, in Problem 2 of Appendix A.

3.2 Ground Superposition

For generating quantifier-free first-order formulas with equalities, over which ground and ordered superposition reasoning had to be employed, we aimed at (i) generating unsatisfiable sets S of ground formulas with uninterpreted functions symbols, such that (ii) refutation proofs of S had similar lengths and complexities. Similar to Section 3.1, we fixed a template for S and only varied the KBO simplification ordering \prec to be used for refuting S within the superposition calculus. To this end, we considered variations of weight functions w and symbol precedence \gg over S , yielding thus different KBOs \prec to be used for refuting S . The main steps of our approach are summarized below.

- (1) We fixed the template for S to be the following set of four clauses

$$E(F(X)) = a \vee E(G(Y)) = a \tag{2}$$

$$F(X) = a [\vee H(b) \neq H(b)] \tag{3}$$

$$G(Y) = a [\vee H(b) \neq H(b)] \tag{4}$$

$$E(a) \neq a [\vee H(b) \neq H(b)], \tag{5}$$

where $E, F, G, H \in \{f, g\}$, $X, Y \in \{a, b\}$, and the literal in $[]$ is added to the clauses optionally.

(2) We created instances of S of this template ensuring that no clause in S is redundant, by considering the following constraints.

- $E \neq H$ and $F(X) \neq G(Y)$;
- Either X or Y is not a . Similarly, either F or G is not E ;
- The literal $H(b) \neq H(b)$ is in exactly one of the clauses (3), (4), (5).

As a result, we produced 12 instances of S satisfying the above properties.

(3) We considered the term algebras induced by the generated instances of S and designed KBO orderings \prec such that refuting the respective instances of S using \prec requires ordering terms both using weight w and precedence \gg . In addition, we imposed that either $F(X) \prec a \prec G(Y)$ or $G(Y) \prec a \prec F(X)$ hold. With such orderings \prec , the shortest refutations of instances of S are of the same length, and in at least one application of superposition, a is replaced by either $F(X)$ or $G(Y)$ in the resulting clause. We generated eight different KBOs \prec fulfilling these conditions, to be used with instances of S , as illustrated in Table 1 of Appendix B.2.

(4) As a result, we obtained 36 different problems (combinations of instances of S and \prec) for the ground superposition reasoning task of our exam; Problem 3 of Appendix A shows such an instance.

4 Implementation

We implemented our approach to randomly generating SAT, SMT, and non-ground first-order problems within Haskell, whereas our ground superposition problem generator was implemented in Python. All together, our toolchain involved about 2300 lines of code, including additional scripts for putting parts together. We encoded each randomly generated SMT and first-order formula within the SMT-LIB input format [1] and, for sanity checks, run the SMT solver Z3 [9] and the first-order theorem prover Vampire [8] for proving the respective formulas. In addition, each formula has been converted to L^AT_EX, yielding randomly generated exam sheets – one such exam sheet is given in Appendix A.

5 Conclusion

We describe a randomized approach and toolchain for generating exam problems in automated reasoning, in particular in the setting of SAT, SMT, and first-order theorem proving. Our approach was used to generate individual exam sheets focused on problem solving within automated deduction, and could be adapted to other constraints and course frameworks.

References

1. Barrett, C., Fontaine, P., Tinelli, C.: The SMT-LIB standard: Version 2.6. Tech. rep., Department of Computer Science, The University of Iowa (2017), available at www.SMT-LIB.org
2. Biere, A., Heule, M., van Maaren, H., Walsh, T. (eds.): Handbook of Satisfiability, Frontiers in Artificial Intelligence and Applications, vol. 185. IOS Press (2009)
3. Claessen, K., Hughes, J.: Quickcheck: a lightweight tool for random testing of haskell programs. In: Proc. ICFP. pp. 268–279 (2000)
4. Davis, M., Logemann, G., Loveland, D.W.: A Machine Program for Theorem-Proving. Commun. ACM **5**(7), 394–397 (1962)
5. Ganzinger, H., Hagen, G., Nieuwenhuis, R., Oliveras, A., Tinelli, C.: DPLL(T): Fast Decision Procedures. In: Proc. of CAV. pp. 175–188 (2004)
6. Gleiss, B., Kovács, L., Rath, J.: Subsumption Demodulation in First-Order Theorem Proving. In: Proc. IJCAR (2020), to appear
7. Gulwani, S., Radicek, I., Zuleger, F.: Automated Clustering and Program Repair for Introductory Programming Assignments. In: Proc. PLDI. pp. 465–480 (2018)
8. Kovács, L., Voronkov, A.: First-Order Theorem Proving and Vampire. In: Proc. CAV. pp. 1–35 (2013)
9. de Moura, L., Bjørner, N.: Z3: an efficient SMT solver. In: Proc. TACAS. pp. 337–340 (2008)
10. Nelson, G., Oppen, D.C.: Simplification by Cooperating Decision Procedures. ACM Trans. Program. Lang. Syst. **1**(2), 245–257 (1979)
11. Nieuwenhuis, R., Rubio, A.: Paramodulation-Based Theorem Proving. In: Handbook of Automated Reasoning, pp. 371–443 (2001)
12. Tinelli, C.: A DPLL-Based Calculus for Ground Satisfiability Modulo Theories. In: Proc. JELIA. pp. 308–319 (2002)
13. Tseytin, G.S.: On the Complexity of Derivation in Propositional Calculus, chap. Studies in Constructive Mathematics and Mathematical Logic, pp. 115–1125. Steklov Mathematical Institute (1970)
14. Wang, K., Singh, R., Su, Z.: Search, Align, and Repair: Data-Driven Feedback Generation for Introductory Programming Exercises. In: Proc. PLDI. pp. 481–495 (2018)

A Appendix A - Randomly Generated Exam Sheet of Automated Deduction

Automated Deduction – SS 2020
Final Exam – June 17, 2020

Version 2020-06-17 / 36

Problem 1. (25 points) Consider the formula:

$$(r \wedge \neg(q \rightarrow p)) \vee (q \leftrightarrow \neg(p \rightarrow q))$$

- (a) Which atoms are pure in the above formula?
- (b) Compute a clausal normal form C of the above formula by applying the CNF transformation algorithm with naming and optimization based on polarities of subformulas;
- (c) Decide the satisfiability of the computed CNF formula C by applying the DPLL method to C . If C is satisfiable, give an interpretation which satisfies it.

Problem 2. (25 points) Consider the formula:

$$\begin{aligned} b = c \wedge f(b+1) \neq b+2 \wedge \text{read}(A, f(c+1)) = c \\ \wedge (\text{read}(A, f(b+1)) = b+3 \vee \text{read}(\text{write}(A, b+2, f(c)), f(c+1)) = c+2) \end{aligned}$$

where b, c are constants, f is a unary function symbol, A is an array constant, $\text{read}, \text{write}$ are interpreted in the array theory, and $+, -, 1, 2, 3, \dots$ are interpreted in the standard way over the integers.

Use the Nelson-Oppen decision procedure in conjunction with DPLL-based reasoning in the combination of the theories of arrays, uninterpreted functions, and linear integer arithmetic. Use the decision procedures for the theory of arrays and the theory of uninterpreted functions and use simple mathematical reasoning for deriving new equalities among the constants in the theory of linear integer arithmetic. If the formula is satisfiable, give an interpretation that satisfies the formula.

Problem 3. (25 points) Consider the KBO ordering \succ generated by the precedence $f \gg a \gg b \gg g$ and the weight function w with $w(f) = 0, w(b) = 1, w(g) = 1, w(a) = 3$. Let σ be a well-behaved selection function wrt \succ . Consider the set S of ground formulas:

$$\begin{aligned} f(g(b)) &= a \vee f(g(a)) = a \\ g(b) &= a \\ g(a) &= a \\ g(b) \neq g(b) \vee f(a) &\neq a \end{aligned}$$

Show that S is unsatisfiable by applying saturation on S using an inference process based on the ground superposition calculus $\text{Sup}_{\succ, \sigma}$ (with the inference rules of binary resolution BR_σ included). Give details on what literals are selected and which terms are maximal.

Problem 4. (25 points) Consider the following inference:

$$\frac{P(h(g(d, d), b)) \vee \neg P(h(f(d))) \vee f(d) \neq h(g(a, a)) \quad \neg P(h(g(x, b))) \vee f(d) \neq h(g(y, y))}{\neg P(h(f(d))) \vee f(d) \neq h(g(a, a))}$$

in the non-ground superposition inference system Sup (including the rules of the non-ground binary resolution inference system BR), where P is a predicate symbol, f, g, h are function symbols, a, b, d are constants, and x, y are variables.

- (a) Prove that the above inference is a sound inference of Sup .
- (b) Is the above inference a simplifying inference of Sup ? Justify your answer based on conditions of clauses being redundant.

B Appendix B

B.1 Details on Random Problem Generation

As we alluded to in Section 2.1, some issues may arise if the restrictions on the randomly generated formula are too strict:

- The sample space might be empty or very sparse. In practice, this manifests as the generator seemingly getting stuck, usually resulting in the process being killed by the user. For example, consider the restriction on polarities of propositional variables. Combined with the other restrictions, it is impossible to get a formula that contains atomic proposition of purely positive and purely negative polarity at the same time.
- The second issue manifests less drastically but is perhaps more problematic: the sample space may be too uniform, leading to the generation of trivial and/or very similar formulas. In particular, we encountered this problem when we restricted the number of models to exactly one, or zero. We note that there simply are not that many ways to rule out 8 interpretations using only 7 connectives.

Regarding the filtering of generated formulas using the constraints discussed in Section 2, we did not require very efficient algorithms since the formulas under consideration are very small. For example, for the restriction on the number of models we used a naive satisfiability test based on evaluating the formula under each possible interpretation. An advantage of this is that the addition of new filters/constraints is relatively easy.

For the random problem generation setting of Section 2, we used design principles from the Haskell library QuickCheck [3]. However, because of our many filtering criteria, we wanted the generator to support backtracking. To this end, we created a simple mtl-style typeclass `MonadChoose` with a single primitive operation `choose` for choosing an element from a list of possible choices:

```
class MonadPlus m => MonadChoose m where
  -- Choose element with uniform probability
  choose :: [a] -> m a
```

Our generator implementations are generic over the monad, constrained by `MonadChoose`. The following listing shows (a slightly simplified) part of the inference generator discussed in Section 2.2.

```

genExamInference :: MonadChoose m => m Inference
genExamInference = do
  -- Define signature (partially omitted)
  let vars = ["x", "y", "z"]
  let opts = GenOptions{ vars = vars, ... }

  -- Choose variables to appear in l1 and l2
  v1 <- choose vars
  v2 <- choose (filter (/= v1) vars)

  -- Generate literals
  -- l1: exactly one occurrence of v1
  l1 <- mfilter ((==1) . length . toListOf variables)
    $ genUninterpretedLiteral opts{ vars = [v1] }
  -- l2: at least two occurrences of v2
  l2 <- mfilter ((>=2) . length . toListOf variables)
    $ genEqualityLiteral opts{ vars = [v2] }
  -- l3: ground literal
  l3 <- genUninterpretedLiteral opts{ vars = [] }

  -- (rest omitted)
  return inference

genEqualityLiteral, genUninterpretedLiteral
  :: MonadChoose m => GenOptions -> m Literal
-- (literal generators omitted)

```

We used two concrete implementations to evaluate generators:

1. `RandomChoice`, a monad that implements `choose` as uniform random choice with backtracking support. Conceptually, this is like the standard list monad where `choose` works like the regular monadic `bind` for lists except that it shuffles the list with a random permutation first. This evaluation method is used to generate random exams.
2. The standard list monad to enumerate the sample space. This second evaluation method helps verifying that the sample space is sufficiently large.

B.2 Weights and Precedences for Ground Superposition

The weights and precedences used to generate the KBOs for the superposition problem from Section 3.2 are displayed in Table 1.

The upper part of the table shows all weight and precedence combinations, denoted as $w_{i,I}, p_{i,I}$ for $i \in \{1, 2, 3, 4\}$ and $I \in \{f, g\}$ (for convenience, the table contains both $w_{i,f}$ and $w_{i,g}$, as well as $p_{i,f}$ and $p_{i,g}$ for all values of i). The lower part of the table displays the values of $i_1, I_1; i_2, I_2$; and i_3, I_3 , corresponding to the three weight and precedence combinations selected for each instance of the clause set.

weight of: $f \ g \ a \ b$	precedence	weight of: $f \ g \ a \ b$	precedence
$w_{1,f} : 1 \ 3 \ 2 \ 1$	$p_{1,f} : a \gg b \gg f \gg g$	$w_{1,g} : 3 \ 1 \ 2 \ 1$	$p_{1,g} : a \gg b \gg g \gg f$
$w_{2,f} : 0 \ 3 \ 2 \ 1$	$p_{2,f} : f \gg a \gg g \gg b$	$w_{2,g} : 3 \ 0 \ 2 \ 1$	$p_{2,g} : g \gg a \gg f \gg b$
$w_{3,f} : 0 \ 1 \ 3 \ 1$	$p_{3,f} : f \gg a \gg b \gg g$	$w_{3,g} : 1 \ 0 \ 3 \ 1$	$p_{3,g} : g \gg a \gg b \gg f$
$w_{4,f} : 1 \ 2 \ 3 \ 1$	$p_{4,f} : g \gg f \gg a \gg b$	$w_{4,g} : 2 \ 1 \ 3 \ 1$	$p_{4,g} : f \gg g \gg a \gg b$
condition		i_1, I_1	$i_2, I_2 \quad i_3, I_3$
$F \neq G$ and $X \neq Y$		1, E	2, E 3, E
$F \neq G$ and $X = Y$		1, H	2, E 4, H
$F = G$ and $X \neq Y$		1, H	2, H 3, E

Table 1. Weights and precedences for the ground superposition problem.